

Stochastic Grammars

Stochastic Context-Free Grammars

I. Holmes

Department of Bioengineering
University of California, Berkeley

Spring semester

Outline

- 1 Overview of transformational grammars
- 2 RNA structure
- 3 Dynamic programming algorithms for SCFGs
- 4 Beyond SCFGs

- Overview: HMM profiles, HMM genefinders, SCFGs for RNA, repeats, beta-sheets; Natural Language Processing
- What is a transformational grammar? Formal definition: terminals Ω , nonterminals Φ , transformation rules
 - “Language” = set of strings generated by the grammar
 - “Parser” = computer program to **decide** if a given input string is in the language (returns “true” or “false”)
 - More generally, we’re interested in parsers that compute scores (energies, probabilities) for a given input string
 - These scores are associated with the transformation rules. The grammar is said to be *score-attributed* (Knuth)

- The Chomsky hierarchy of grammars and their associated parsers
 - Regular grammars: finite-state machines (HMMs)
 - Context-free grammars: pushdown automata (SCFGs)
 - Context-sensitive grammars: finite-tape Turing machines (linear-bounded automata)
 - Unrestricted grammars: infinite-tape Turing machines

- The Chomsky hierarchy of grammars and their associated parsers (lower levels)
 - Regular grammars: finite-state machines (HMMs)
 - Context-free grammars: pushdown automata (SCFGs)
 - The **parse tree**; the **inside sequence** and **outside sequence**
 - Chomsky Normal Form; Eddy *et al*'s "RNA Normal Form"
 - Transformations to & consequent universality of Chomsky Normal Form
 - History: Panini's "Ashtadhyayi": a CFG for Sanskrit (3959 rules (sutras), circa.300-700BC). Bishop Robert Lowth's "A Short Introduction to English Grammar" (1762). Chomsky's theory of universal generative grammars (1956 and onwards).

- The Chomsky hierarchy of grammars and their associated parsers (upper levels)
 - Context-sensitive grammars: Turing machines with bounded tape = linear-bounded automata
 - Big category: low-complexity sequence repeats, tandem repeats, other bounded correlations e.g. pseudoknots
 - Lempel-Ziv compression algorithm (allowing local stutter) can be viewed as one of these
 - Limited context sensitivity (e.g. basepair stacking) can be effected using a *lexicalised* SCFG
 - Tree-Adjoining Grammars, Linear-Indexed Grammars, etc.
 - Unrestricted grammars: complete Turing machines

- Summary of Chomsky hierarchy (NB $\text{regular} \subset \text{context-free} \subset \text{context-sensitive} \subset \text{unrestricted}$)

Class	Example rule	Automaton	Complexity
Regular	$S \rightarrow x S$	Finite-state	$O(L)$
Context-free	$S \rightarrow T U$	Pushdown	$O(L^3)$
Context-sensitive	$S T \rightarrow TS$	Linear-bounded	$O(\Phi ^L)$
Unrestricted	$S T \rightarrow U$	Turing machine	Undecidable

Here S, T, U are nonterminals and x is a terminal.

“Complexity” refers to the time complexity of parsing a sequence of length L .

- Why RNA is important in evolution and cell biology
 - RNA world; pre- and post-transcriptional regulation; Crick's idea of studying simple examples; ribotechnology

- RNA structure terminology: basepairs, stems, loops, pseudoknots, kissing loops

Terms contributing to the free energy of a folded RNA structure (scor.berkeley.edu)

- Hydrogen bonding between bases. Canonical, noncanonical pairs
- Stacking energies due to overlap of π -orbitals of adjacent planar basepairs
 - H-bonding and stacking terms can be combined and measured by direct experiment.
- Unusual configurations: tetraloops, triloops, triple-A platforms
 - Finite number of cases, so also amenable to experimental measurement.

- Entropic cost of closing loops
 - Theory: rods and Gaussian springs. Integrate out displacements, get likelihood ratio (Doi-Edwards, Isambert)
 - Statistics of random walk, $\langle |\Delta \mathbf{x}|^2 \rangle \propto t$, and self-avoiding walk, $\langle |\Delta \mathbf{x}|^2 \rangle \propto t^{1+\epsilon}$
 - Renormalisation (Edwards, de Gennes)
 - Empirical scaling laws fit experimental measurements
- Ligands: solvation, metal ions, small molecules
 - General rules not yet known. Specific small-molecule binding requires conserved motifs (riboswitches).
- Unlike proteins (where amino acid sidechains make multiple contacts), a “**basepair stacking + convex loop penalty**” picture of the free energy seems reasonably accurate as a first approximation (Zuker, Turner, Mathews...)

- The Nussinov algorithm: Finds *strictly nested* foldback structures, i.e. excluding pseudoknots.
 - RNA sequence X , length L , nucleotides $x_1 \dots x_L$
 - Let $H(x, x') = 1$ if xx' is a canonical Watson-Crick basepair, and 0 otherwise
 - Nussinov recursion finds the structure for $x_i \dots x_j$ that has the most strictly nested canonical basepairs

$$S(i, j) = \max \left(\begin{array}{c} S(i+1, j), \\ S(i, j-1), \\ S(i+1, j-1) + H(x_i, x_j), \\ \max_{i \leq k \leq j} (S(i, k) + S(k, j)) \end{array} \right)$$

Best structure for X is found by traceback from $S(1, L)$

Nussinov algorithm = parser for **score-attributed grammar**

Rule		Score
S	$\rightarrow S x$	0
	$ x S$	0
	$ x S x'$	$H(x, x')$
	$ S S$	0
	$ \epsilon$	0

If we allow $H(x, x')$ to be the free energy of basepair formation for xx' (and assume a zero energy cost for any structural feature except basepairs), then this becomes an **energy-attributed grammar**. Same recursion for S now calculates “ground state” energy for a very simple energy model ignoring all but hydrogen bonding terms for strictly nested basepairs.

The partition function for this highly simplified model is

$$Z(i,j) = Z(i+1,j) + Z(i,j-1) + Z(i+1,j-1) \exp(-\beta H(x_i, x_j)) \\ + \sum_{i \leq k \leq j} Z(i,k)Z(k,j)$$

where $\beta = 1/kT$ is an “inverse temperature”. NB: have simply changed $H \rightarrow \exp(-\beta H)$, $+$ $\rightarrow \times$ and $\max \rightarrow +$ in equation for S

We can also have **probability-attributed grammars** or **stochastic grammars**. Relationship between scoring schemes:

Score	Form
Bayes	$P(x)$
Shannon	$h(x) = -\log_2 P(x)$
Boltzmann	$E(x) \sim -\log_e P(x)$

Specifically the Boltzmann probability uses a scaling factor (inverse temperature) and a partition function,

$$P(x) = \frac{1}{Z} \exp[-\beta E(x)], \text{ where } Z = \sum_x \exp[-\beta E(x)].$$

Energy-attributed grammar (Zuker, MFOLD). Parse tree for subsequence $X_i \dots X_j$ must be rooted at W or $V_{X_i X_j}$.

LHS		RHS	Energy
W	\rightarrow	$W x$	0
	$ $	$x W$	0
	$ $	$x V_{xy} y$	$\alpha(x, y)$
	$ $	$W W$	0
	$ $	$x y$	$\alpha(x, y)$
	$ $	ϵ	0
V_{ab}	\rightarrow	z^n	$h(n)$
	$ $	$x V_{xy} y$	$\alpha_S(x, y a, b)$
	$ $	$x V_{xy} y z^n$	$\alpha(x, y) + b(n)$
	$ $	$z^n x V_{xy} y$	$\alpha(x, y) + b(n)$
	$ $	$z^m x V_{xy} y z^n$	$\alpha(x, y) + i(m + n)$
	$ $	$x V_{xy} y W x' V_{x'y'} y'$	$\alpha(x, y) + \alpha(x', y')$
	$ $	ϵ	0

Here x, y, x', y', z_i are terminals (nucleotides), and z^n is shorthand for an n -nucleotide emission $z_1 \dots z_n$.

The scoring scheme is as follows:

Free energy term	Meaning
$\alpha(x, y)$	Energy of basepair xy (end of stem, no stacking)
$\alpha_S(x, y a, b)$	Energy of basepair xy stacked on top of basepair ab
$h(n)$	Hairpin loop enclosing n bases
$b(n)$	Asymmetric bulge of n bases
$i(n)$	Interior loop (symmetric bulge) with total of n bases

This scoring scheme was historically formulated in terms of Sankoff's " k -loop decomposition".

NB Zuker's algorithm (grammar), like Nussinov's, excludes pseudoknots.

Strictly, Zuker described the algorithm that finds the lowest-energy parse using the above grammar (CYK).

McCaskill described the algorithm for calculating the partition function (Inside) and thus the posterior probabilities of individual basepairs (Outside).

- Chomsky normal form. (RNA normal form is more useful in practise, but CNF is easier to present.)

For nonterminals $A, B, C \in \Phi$ and terminals $a \in \Omega$:

Rule	Name
$A \rightarrow BC$	Bifurcation
$A \rightarrow a$	Emission
$A \rightarrow \epsilon$	Termination

Probabilities denoted by $P(\text{rule})$, e.g. $P(A \rightarrow BC)$

- Inside algorithm.

Let $I_A(i, k) = P(x_i \dots x_{i+k} | A)$ be sum of probabilities for parse trees rooted in A generating sequence $x_i \dots x_{i+k}$.

$$I_A(i, k) = \left(\sum_B \sum_C \sum_{j=0}^k P(A \rightarrow BC) I_B(i, j) I_C(i + j, k - j) \right) + \left\{ \begin{array}{ll} 0 & \text{if } k > 1 \\ P(A \rightarrow x_i) & \text{if } k = 1 \\ P(A \rightarrow \epsilon) & \text{if } k = 0 \end{array} \right\}$$

NB loopy dependencies, e.g. if $P(A \rightarrow AA) \neq 0$ and $P(A \rightarrow \epsilon) \neq 0$. It's common to try to avoid these when designing the grammar.

- Cocke-Younger-Kasami (CYK) algorithm.

Let $Y_A(i, k)$ be probability of ML parse tree rooted in A generating sequence $x_i \dots x_{i+k}$.

$$Y_A(i, k) = \max \left(\begin{array}{l} \max_B \max_C \max_{j=0}^k P(A \rightarrow BC) Y_B(i, j) Y_C(i + j, k - j), \\ \left\{ \begin{array}{ll} 0 & \text{if } k > 1 \\ P(A \rightarrow x_i) & \text{if } k = 1 \\ P(A \rightarrow \epsilon) & \text{if } k = 0 \end{array} \right\} \end{array} \right)$$

NB similar to Nussinov.

- Outside algorithm.

Let $O_A(i, k) = P(x_1 \dots x_{i-1}, A, x_{i+k+1} \dots x_L | S)$ be sum of probabilities for incomplete parse trees rooted in S , ending in A and generating outside sequence $x_1 \dots x_{i-1}$ and $x_{i+k+1} \dots x_L$.

$$O_A(i, k) = \sum_B \sum_C \left(\sum_{j=0}^i O_C(i-j, j+k) P(C \rightarrow BA) I_B(i-j, j) + \sum_{j=0}^{L-i-k} O_C(i, j+k) P(C \rightarrow AB) I_B(i+k, j) \right)$$

with boundary condition $O_A(0, L) = \delta(A = S)$.

- Inside-Outside and posterior probabilities.

Posterior probability that parse tree contains a subtree with inside sequence $x_i \dots x_{i+k}$ rooted in state A :

$$P(A_{i,i+k} | X) = \frac{O_A(i, k) I_A(i, k)}{I_S(0, L)}$$

Posterior probability of bifurcation $A_{i,j+k} \rightarrow B_{i,j} C_{i+j,k}$:

$$P(A_{i,j+k} \rightarrow B_{i,j} C_{i+j,k} | X) = \frac{O_A(i, j+k) P(A \rightarrow BC) I_B(i, j) I_C(i+j, k)}{I_S(0, L)}$$

etc.

- Parameter estimation: we can write the parse tree likelihood as $\prod_i \theta_i^{n_i}$, where θ_i is the probability of rule i and n_i is the number of times it was applied. As with HMMs, the EM algorithm for SCFGs thus involves computing posterior expectations $\langle n_i \rangle$ for these counts, and setting $\theta_i \propto \langle n_i \rangle$. The posterior expectations are computed using the Inside-Outside algorithm, as shown for rules of the form $P(A \rightarrow BC)$.
- There is a KYC-like analogue to Outside, that can be used to find ML parse tree including a particular subsequence.
- Implementation issues: time complexity is $O(|\Phi|^3 L^3)$, memory is $O(|\Phi| L^2)$.
 - Cubic factors in time complexity arise due to bifurcations, so minimize number of bifurcations for max efficiency (c.f. RNA normal form).

Pair SCFGs, evolutionary SCFGs and tree transducers

- Evolutionary SCFGs: PFOLD (xfold, evofold, etc.)
 - As with Evolutionary HMMs, we can let the terminals be alignment columns
 - Again, terminal emission likelihood $P(A \rightarrow a)$ is implemented as Felsenstein pruning

- Pair SCFGs: Evoldoer. Version of TKF that describes evolution of RNA secondary structure (Holmes 2005).
 - Two kinds of TKF91 links model, recursively nested in a tree
 - Stem sequences rooted in S nonterminals; basepair alphabet Ω^2 ; ends in an L
 - Loop sequences rooted in L nonterminals; nucleotide alphabet Ω ; S 's also allowed in sequence
 - Really need to adapt TKF91 model to allow deletion prob to depend on symbol, otherwise S substructures get deleted at same rate as nucleotides
- Pair SCFGs (e.g. Stemloc, QRNA). Heuristic, but with lots of go-faster stripes (pre-aligning & pre-folding sequences).
- Rules for composing Pair SCFGs exist (c.f. string transducers; can view conditionally normalized pair SCFGs as “tree transducers”).

Graph grammars

- Tree-adjoining grammars
 - Aravind Joshi (1975, 1985)
- Rivas-Eddy papers
 - A dynamic programming algorithm for RNA structure prediction including pseudoknots. JMB 1999.
 - The language of RNA: a formal grammar that includes pseudoknots. Bioinformatics 2000.
- Graph grammars: easy to describe and simulate, attractive for biology; but how does their DP work?
- Other grammars whose marginals are easy to compute by sum-product DP, e.g.
 - stochastic tree grammars (Abe and Mamitsuka, ISMB 1994)
 - context-sensitive HMMs with finite memory of last N emitted characters (Yoon and Vaidyanathan, 2004)

Discriminative grammars

Discriminative grammars: conditional log-linear models

- Recall HMMs and linear CRFs form a “generative-discriminative pair”
- The analogous discriminative model for SCFGs is a Conditional Log-Linear Model
- This allows a great deal of physics-like parameterization (loop entropies, stacking free energies, terminal mismatch...) without having to introduce many new nonterminals
 - Do, Woods and Batzoglou. “CONTRAFold: RNA secondary structure prediction without physics-based models.” Bioinformatics 22:14, pp e90-e98

Summary

- SCFGs